

Notebook3

March 17, 2018

Contents

1	Comparing the Convolution to the Transform	1
2	Implementation of the test	1
3	Results	3

1 Comparing the Convolution to the Transform

- Section 3.4 in the book describes that due to how the math works out multiplication in Fourier space $F(\omega)H(\omega) = f(x) * h(x)$ where $F(\omega)$ and $H(\omega)$ are the Fourier transform on $f(x)$ and $h(x)$ respectively.
- For my tests I have also compared the convolution to multiplication in Cosine Transform space, which lead to interesting results.

2 Implementation of the test

- Since I did tests with both the cosine and fourier transformers, Notebook2 has the filters for dct which won't be reposted.
- At first I wasn't sure how to convert a Double into a \mathbb{C} number, so I only tested the dct first, however I later figured out that I can just set the imaginary number to 0 for images.

```
repaFft :: Shape sh => Array V sh Double -> Array V sh (C Double)
repaFft = repaVecComp (fft . fmap (:+ 0))
```

```
repaIFft :: Shape sh => Array V sh (C Double) -> Array V sh Double
repaIFft = repaVecComp (fmap magnitude . ifft)
```

```
repaFftP :: (Monad m, Load r sh Double) => Array r sh Double -> m (Array V sh (C Double))
repaFftP = fmap repaFft . computeVectorP
```

```
repaIFftP :: (Monad m, Load r sh (C Double)) => Array r sh (C Double) -> m (Array V sh Double)
repaIFftP = fmap repaIFft . computeVectorP
```

- When converting an image to fft, we just add 0 to the imaginary part, and when we compute the inverse, we just take the magnitude, safely converting the image back into a Double from being Complex.
- Besides the above function for FFT transforms, I noticed that after running my tests the image was off by around 3 pixels (not 2 but not quite 3 actually), so the below function was made

```
offsetFft :: Source r b => Array r DIM2 b -> Array D DIM2 b
```

```
offsetFft arr = R.traverse arr id f
```

```
  where sh@(Z :: i :: j) = extent arr
```

```
    f index (Z :: x :: y)
```

```
      | isInside2 sh newShape = index newShape
```

```
      | x + 2 ≥ i ∧ y + 2 ≥ j = index (ix2 (x + 2 - i) (y + 2 - j))
```

```
      | x + 2 ≥ i              = index (ix2 (x + 2 - i) (y + 2))
```

```
      | otherwise              = index (ix2 (x + 2)      (y + 2 - j))
```

```
    where newShape = ix2 (x + 2) (y + 2)
```

- This code just creates a new array where all the pixels are shift down and right by two and wrap around when encountering an edge.

- this code can easily be made more general instead of shifting everything by 2 you pass in a pad down and a pad up.

- Now that we have all the transformers we need, we can start to make our filters and padding our filters.

```
-- the one given in the python code is wrong, as it uses [1,2,6,2,1]
```

```
-- it seems even the lecture is wrong, because if yo add them all up, they
```

```
-- don't add up to 256.. can check by summing over my gaussian and seeing its 256
```

```
gaussian :: Array U DIM2 Double
```

```
gaussian = fromListUnboxed (ix2 5 5) $ (*) . (/ 256) <$> [1,4,6,4,1] <*> [1,4,6,4,1]
```

- Here we create the 5 by 5 Gaussian filter array by using the generalized cross product trick from Notebook1, then transforming the list into an array

```
pad :: (Source r e) => e -> DIM2 -> Array r DIM2 e -> Array D DIM2 e
```

```
pad val sh vec = fromFunction sh makePad
```

```
  where
```

```
    Z :: i :: j = R.extent vec
```

```
    makePad sh@(Z :: x :: y)
```

```
      | x ≥ i ∨ y ≥ j = val
```

```
      | otherwise     = vec ! sh
```

```
padOff :: (Source r e) => e -> DIM2 -> Array r DIM2 e -> Int -> Int -> Array D DIM2 e
```

```
padOff val sh vec offx offy = fromFunction sh makePad
```

```
  where
```

```
    Z :: i :: j = R.extent vec
```

```
    makePad sh@(Z :: x :: y)
```

```
      | x - offx ≥ i ∨ x - offx < 0
```

```
      ∨ y - offy ≥ j ∨ y - offy < 0 = val
```

```
      | otherwise              = vec ! ix2 (x - offx) (y - offy)
```

- While creating the testDiffGen which will be displayed after meanDiff, I noticed that I don't have any padding function that would allow me to make the Gaussian the same size as our image

- There are two pad's, padOff allows the user to specify where they want their original vector to be in the padded vector.

```
meanDiff :: (Source r c, Fractional c, Source r2 c) => Array r DIM2 c -> Array r2 DIM2 c -> c
```

```
meanDiff arr1 = (/ fromIntegral (i * j)) . sumAllS . R.zipWith (\x y -> abs (x - y)) arr1
```

```
  where Z :: i :: j = R.extent arr1
```

- Here we are simply taking two arrays and subtracting every place then sum up the 2d Vector then dividing by the length of the array. Really this is a simple difference calculation

- Now that all the prep work is done we can finally do our calculation

```
testDiffGen forwardTransformP inverseTransformP forwardTransform paddingP path = do
  img      ← readIntoRepa path
  let origV = R.map fromIntegral (repaRGBToGrey img)
  origU     ← computeUnboxedP origV
  convolved ← convolveOutP outClamp gaussian origU
  convolvedC ← forwardTransformP . delay $ convolved

  cosOrigV ← forwardTransformP origV
  let cosOrigU = computeUnboxedS (delay cosOrigV)
  let ext@(Z :: i :: j) = R.extent cosOrigU
  let padding | paddingP = padOff 0 ext gaussian (i `div` 2) (j `div` 2)
              | otherwise = pad    0 ext gaussian

  paddedGausV ← computeVectorP $ pad 0 (R.extent cosOrigU) gaussian
  paddedGausU ← computeUnboxedP (delay $ forwardTransform paddedGausV)

  let matrixMultC = (cosOrigV *~ paddedGausU)
  -- convert it back with idct
  matrixMult ← inverseTransformP matrixMultC
  saveRepaGrey "test.png" matrixMult
  saveRepaGrey "test2.png" convolved
  print ("difference in the DCT "      <> show (meanDiff matrixMultC convolvedC))
  print ("difference in the NormalPlane " <> show (meanDiff matrixMult convolved))
  return (matrixMult, convolved)

testDiffDct = testDiffGen repaDctImageP repaIDctImageP repaDct

testDiffFft = testDiffGen repaFftP (fmap (computeVectorS . offsetFft) . repaIFftP) repaFft True


-- testDiffGen takes the forward transform and inverse transform, both with parallel computation,
-- a normal forward transform and where padding should be included and a path for the image.
-- In the computation, we first read the image into a Repa array.
-- then we turn it grey in origV.
-- origU is just origv but unboxed.
-- After we set up our vector we can convolve the Gaussian with image and then transform it.
-- The next section calls the forward transformation on the original image and starts to pad the
-- Gaussian, the Gaussian is in the middle if paddingP is true, or in the top left if paddingP is false.
-- The final section just does the matrix multiplication, then calls the inverse and saves the image
-- before printing the difference between the Fourier/Cosine plane and the identity plane.
-- testDiffDct and testDiffFft just call the generalized function with different default functions
-- which have been discussed in this notebook and notebook2.
```

3 Results


- Now that we have all the functions in place, we can run testDiffDct and testDiffFft and see what results we get
- The convolution answer for all of the images seen below is the following



- `y <- testDiffDct False "./data/bunny.png" _`
 - "difference in the DCT 2702.16509984095"
 - "difference in the NormalPlane 39.963790639071824"



- as we can see it just turns the dark portions of the bunny white and is oddly connected. Also note that the image is also somewhat blurred so it seemed the filter somewhat worked
- As expected the difference in both planes is quite high
- a previous iteration where I centered the Gaussian gave me this



- Which implies that multiplication of the DCT doesn't give anywhere near close to being the same as the convolution. Which is expected since we are not using the Fourier Transformer
- `y <- testDiffFft "../data/bunny.png" _`
 - I'm going to post 3 different versions of this, one with 0 bits shifted to the right and down, 2 bits and 3 bits
 - so with two pixels shifted we get



"difference in the DCT 1385.7172620319634 :+ 0.0"
"difference in the NormalPlane 0.1172603668112668"

- We can see that the difference in the normal plane is 0.11
- And we can see that the image is wrapped around a bit on the top and a bit on the left and right but overall the image is the exactly the same
- this type of error would skyrocket the difference which would be much much lower (the given python code was E-32)
- With 3 pixels shifted we get



"difference in the DCT 1385.7172620319634 :+ 0.0"

"difference in the NormalPlane 2.411222994281215"

- And we can see that the image is wrapped around a bit on the far right side and a bit on the bottom
- The difference is greater, which implies that it's closer to being 2 pixels shifted than 3



* This is the image with no offset after the FFT, here we can see clearly how it is shifted.